



Ejercicio 5: Redes Neuronales

Factos

- Estas en la parte pesada de numpy
- Después del ejercicio 5 habrá menos implementaciones de numpy



Preguntas posibles

- P: ¿Necesitamos redes profundas?

R: Sí. Las capas múltiples permiten un mayor poder de abstracción con un presupuesto computacional fijo en comparación con una capa única. Es mejor para la generalización.

- P: Entonces, ¿simplemente construimos redes de 100 capas de profundidad?

R: No es trivial

- Restricciones: Memoria, vanishing gradients, ...
- más profundo != funciona mejor



Ejercicio 5

Recap: Ejercicio 4

Ex4:

- Small dataset
And simple objective
- Simple classifier
Single weight matrix
- Gradient descent solver
Whole forward pass in memory



Ex5:

- CIFAR10
Actual competitive task
- Modularized Network
Chain rule rules
- Stochastic Descent

Recap: Ejercicio 4

```
class Classifier(Network):
    """
    Classifier of the form  $y = \text{sigmoid}(X * W)$ 
    """

    def __init__(self, num_features=2):
        super(Classifier, self).__init__("classifier")

        self.num_features = num_features
        self.W = None

    def initialize_weights(self, weights=None):
        """
        Initialize the weight matrix W

        :param weights: optional weights for initialization
        """
        if weights is not None:
            assert weights.shape == (self.num_features + 1, 1), \
                "weights for initialization are not in the correct shape"
            self.W = weights
        else:
            self.W = 0.001 * np.random.randn(self.num_features + 1, 1)
```

```
def forward(self, X):
    """
    Performs the forward pass of the model.

    :param X: N x D array of training data. Each row is a D-dimensional point.
    :return: Predicted labels for the data in X, shape N x 1
             1-dimensional array of length N with classification scores.
    """
    assert self.W is not None, "weight matrix W is not initialized"
    # add a column of 1s to the data for the bias term
    batch_size, _ = X.shape
    X = np.concatenate((X, np.ones((batch_size, 1))), axis=1)
    # save the samples for the backward pass
    self.cache = X
    # output variable
    y = None

    #####
    # TODO: Implement the forward pass and return the output of the model. Note #
    # that you need to implement the function self.sigmoid() for that #
    #####

    y = X.dot(self.W)
    y = self.sigmoid(y)

    #####
    #                                     END OF YOUR CODE                                     #
    #####
```

Nueva Modularización

Chain Rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial d} \cdot \frac{\partial d}{\partial y}$$



```
class Sigmoid:
    def __init__(self):
        pass

    def forward(self, x):
        """
        :param x: Inputs, of any shape

        :return out: Output, of the same shape as x
        :return cache: Cache, for backward computation, of the same shape as x
        """

    def backward(self, dout, cache):
        """
        :return: dx: the gradient w.r.t. input X, of the same shape as X
        """
```

Resumen: Ejercicio 5

- 1 notebook: pero largo...
- Múltiples implementaciones pequeñas

Definition

$$CE(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N \sum_{k=1}^C \left[-y_{ik} \log(\hat{y}_{ik}) \right]$$

where:

- N is again the number of samples
- C is the number of classes
- \hat{y}_{ik} is the probability that the model assigns for the k 'th class when the i 'th sample is the input.
- $y_{ik} = 1$ iff the true label of the i th sample is k and 0 otherwise. This is called a [one-hot encoding](#).

Task: Check Formula

Check for yourself that when the number of classes C is 2, then binary cross-entropy is actually equivalent to cross-entropy.



Nos vemos el próximo lunes 😊