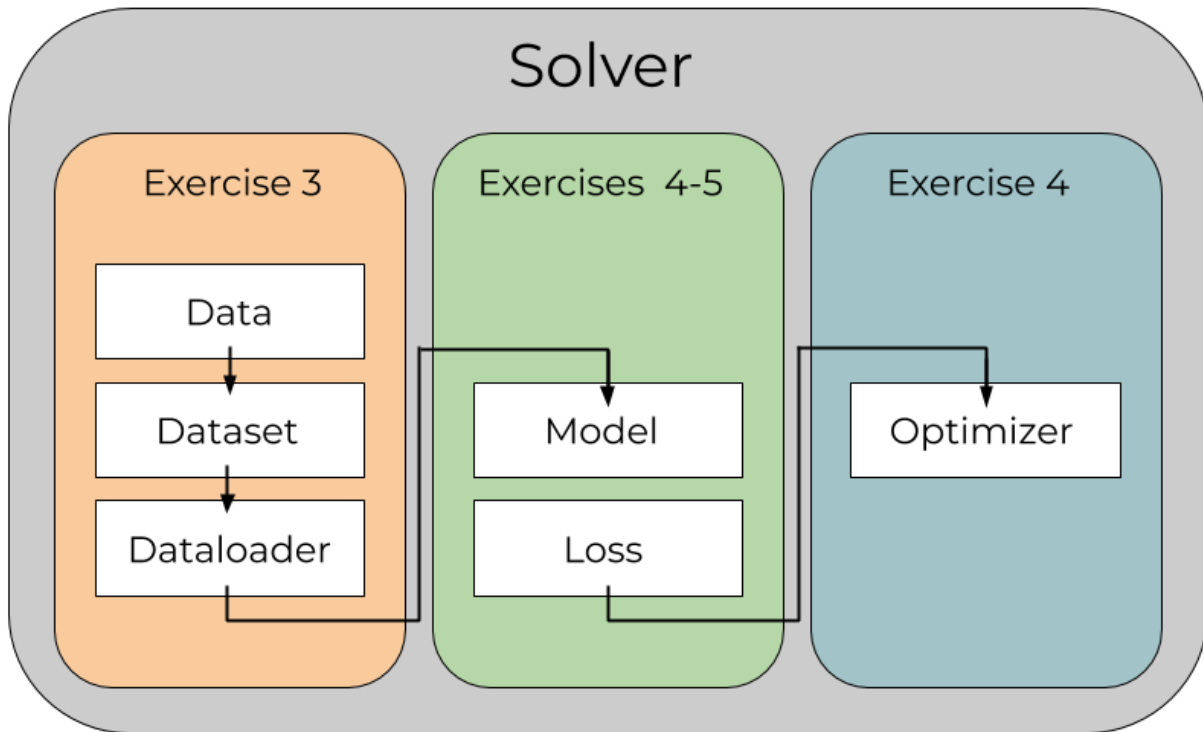


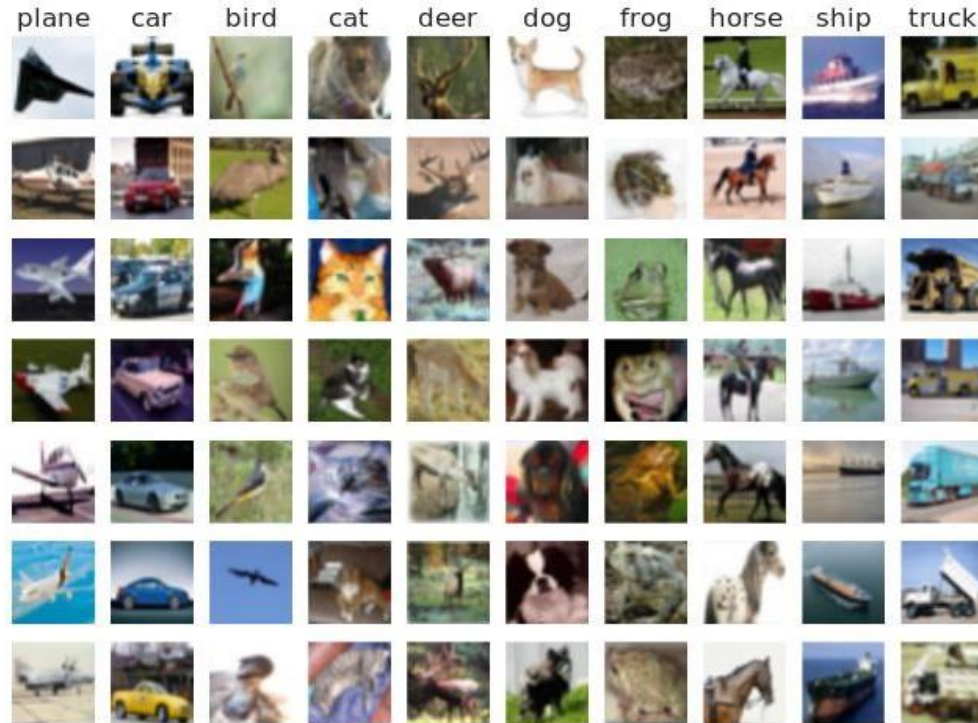
# Ejercicio 6:

# Hyperparameter Tuning

# Repaso: Pillars of Deep Learning



# Objetivo del Ejercicio 6



Cifar10

# Objetivo del Ejercicio 6

- Utilizar las implementaciones existentes
  - Aplicaciones revisadas de ejercicios anteriores
  - Le proporcionaremos implementaciones adicionales de todas las herramientas necesarias para ejecutar los métodos de ejemplo propuestos en la clase.
- Aprenda de las neural network debugging strategies and hyperparameter search



# Leaderboard

- La **precisión** de tu modelo es lo único que cuenta.
  - Al menos un **48%** para aprobar el examen
  - Habrá una **clasificación** de todos los estudiantes.

## Leaderboard

The leaderboard shows for each exercise the highest scoring submission from each user. Only valid submissions are displayed.

Exercise 1	Exercise 3	Exercise 4	Exercise 5	Exercise 6	Exercise 7	Exercise 8	Exercise 9	Exercise 10	Exercise 11
#	User			Score					
1	a0008			100.00					
2	a0001			100.00					
3	a0003			100.00					
4	u0306			100.00					
5	u1540			100.00					

# Previously: Dataset

```
class ImageFolderDataset(Dataset):
    """CIFAR-10 dataset class"""
    def __init__(self, transform=None, mode='train',
                 limit_files=None,
                 split={'train': 0.6, 'val': 0.2, 'test': 0.2},
                 *args, **kwargs): ...

    @staticmethod
    def _find_classes(directory): ...

    def select_split(self, images, labels, mode): ...

    def make_dataset(self, directory, class_to_idx, mode): ...

    def __len__(self): ...

    @staticmethod
    def load_image_as_numpy(image_path): ...

    def __getitem__(self, index): ...
```

```
# Create a train, validation and test dataset.
datasets = {}
for mode in ['train', 'val', 'test']:
    crt_dataset = ImageFolderDataset(
        mode=mode,
        root=cifar_root,
        download_url=download_url,
        transform=compose_transform,
        split={'train': 0.6, 'val': 0.2, 'test': 0.2}
    )
    datasets[mode] = crt_dataset
```

# Previously: Data Loader

```
class DataLoader:
    """
    Dataloader Class
    Defines an iterable batch-sampler over a given dataset
    """
    def __init__(self,
                 dataset,
                 batch_size=1,
                 shuffle=False,
                 drop_last=False): ...

    def __iter__(self): ...

    def __len__(self): ...
```

```
# Create a dataloader for each split.
dataloaders = {}
for mode in ['train', 'val', 'test']:
    crt_dataloader = DataLoader(
        dataset=datasets[mode],
        batch_size=256,
        shuffle=True,
        drop_last=True,
    )
    dataloaders[mode] = crt_dataloader
```

# Previously: Solver

```
class Solver(object):
    """
    A Solver encapsulates all the logic necessary for training classification
    or regression models.
    The Solver performs gradient descent using the given learning rate.
    """

    def __init__(self, model, train_dataloader, val_dataloader,
                 loss_func=CrossEntropyFromLogits(), learning_rate=1e-3,
                 optimizer=Adam, verbose=True, print_every=1,
                 lr_decay = 1.0, **kwargs): """

    def _reset(self): """

    def _step(self, X, y, validation=False): """

    def train(self, epochs=100, patience = None): """

    def get_dataset_accuracy(self, loader): """

    def update_best_loss(self, val_loss, train_loss): """
```

```
solver = Solver(model,
                 dataloaders['train'],
                 dataloaders['val'],
                 learning_rate=0.001,
                 loss_func=MSE(),
                 optimizer=SGD)
```

```
solver.train(epochs=epochs)
```



# Previously: Classification Network

```
class ClassificationNet(Network):  
    """  
    A fully-connected classification neural network with configurable  
    activation function, number of layers, number of classes, hidden size and  
    regularization strength.  
    """  
  
    def __init__(self,  
        activation=Sigmoid(), num_layer=2,  
        input_size=3 * 32 * 32, hidden_size=100,  
        std=1e-3, num_classes=10, reg=0, **kwargs):  
        """  
        """  
  
    def forward(self, X):  
        """  
        """  
  
    def backward(self, dy):  
        """  
        """  
  
    def save_model(self):  
        """  
        """  
  
    def get_dataset_prediction(self, loader):  
        """  
        """
```

```
# Instantiate a new model.  
model = ClassificationNet(activation=Sigmoid(),  
                          num_layer=num_layer,  
                          reg=reg,  
                          num_classes=10)  
  
# X is a batch of training features  
# X.shape = (batch_size, features_size)  
y_out = model.forward(X)  
  
# dout is the gradient of the loss function  
# w.r.t the output of the network.  
# dout.shape = (batch_size, )  
model.backward(dout)
```

# Previously: Binary Cross Entropy Loss

$$BCE(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N \left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right]$$

Where

- $N$  is the number of samples
- $\hat{y}_i$  is the network's prediction for sample  $i$
- $y_i$  is the ground truth label (0 or 1)

# New: Multiclass Cross Entropy Loss

$$CE(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N \sum_{k=1}^C \left[ -y_{ik} \log(\hat{y}_{ik}) \right]$$


Donde

- $N$  es el número de muestras
- $\hat{y}_{ik}$  es la probabilidad predicha por la red para la  $k$ -ésima class dada la muestra  $i$
- $y_{ik}$  es la etiqueta de verdad que es 1 si la muestra  $i$  es de la clase  $k$  o cero en caso contrario.

Implementamos esto por ustedes

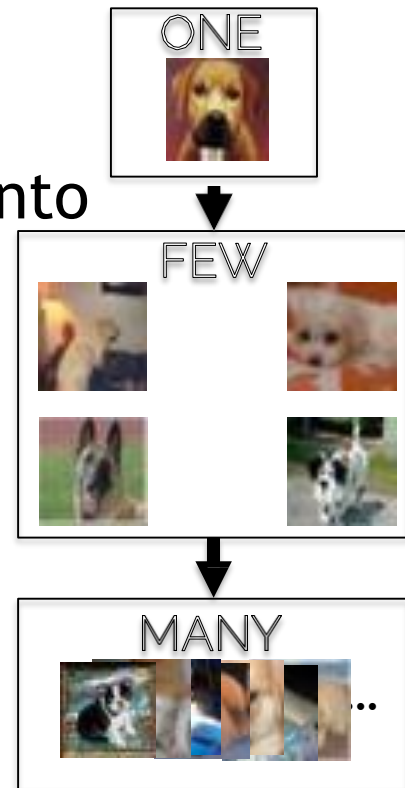
# Receta básica para el ML

- Split your data




# Como empezar

- Empezar con una sola muestra de entrenamiento
  - Comprobar si la salida es correcta
  - Overfit -> La precisión debe ser del 100% porque la entrada acaba de ser memorizada
- Incrementar a un puñado de muestras
- Ir del overfitting a más muestras
  - En algún punto, tendrían que ver generalización




# Como empezar

- Overfit a single training sample



- Then a few samples




# Hyperparameters

- Network architecture (e.g., num layers, hidden layer, activation function)
- Number of iterations
- Learning rate(s) (i.e., solver parameters, decay, etc.)
- Regularization (more later next lecture)
- Batch size
- ...

# Hyperparameter Tuning



**Hyperparameters**



**Parameters**



**Score**



Source: <https://images.deepai.org/glossary-terms/05c646fe1676490aa0b8cab0732a02b2/hyperparams.png>



# ¿Cómo encontrar buenos Hyperparameters?

- Manual Search (trial and error)
- Automated Search:
  - Grid Search
  - Random Search
- Piensen cómo diferentes hyperparametros pueden afectar al modelo
  - E.g. Overfitting? -> Increase Regularization Strength, decrease model capacity

```
from exercise_code.hyperparameter_tuning import grid_search

best_model, results = grid_search(
    dataloaders['train_small'], dataloaders['val_500files'],
    grid_search_spaces = {
        "learning_rate": [1e-2, 1e-3, 1e-4, 1e-5, 1e-6],
        "reg": [1e-4, 1e-5, 1e-6]
    },
    epochs=10, patience=5,
    model_class=ClassificationNet)
```

# Plan de la Práctica: Recap y Outlook

Exercise 03: Dataset and Dataloader  
Exercise 04: Solver and Linear Regression  
Exercise 05: Neural Networks  
Exercise 06: Hyperparameter Tuning

**Numpy**  
**(Reinvent the wheel)**

Exercise 07: Introduction to Pytorch  
Exercise 08: MNIST with Pytorch

**Pytorch/Tensorboard**

Exercise 09: Convolutional Neural  
Networks  
Exercise 10: Semantic Segmentation  
Exercise 11: Recurrent Neural Networks

**Applications**  
**(Hands-off)**

**Nos vemos el  
próximo lunes 😊**