

Inteligencia Artificial

Ejercicio 7: Pytorch

Pytorch

Cronograma de Ej

Exercise 01: Organization
Exercise 02: Math Recap

Intro

Exercise 03: Dataset and Dataloader
Exercise 04: Solver and Linear Regression
Exercise 05: Neural Networks
Exercise 06: Hyperparameter Tuning

Numpy
(Reinvent the wheel)

Exercise 07: Introduction to Pytorch
Exercise 08: Autoencoder

Pytorch/Tensorboard

Exercise 09: Convolutional Neural
Networks
Exercise 10: Semantic Segmentation
Exercise 11: Recurrent Neural Networks

Applications
(Hands-off)

Deep Learning Frameworks

Dos grandes


- Tensorflow - Google
 - As well as Keras
- Pytorch - Facebook

Otros ejemplos

- CNTK - Microsoft
- Mxnet - Apache
- Jax - Google
- ...

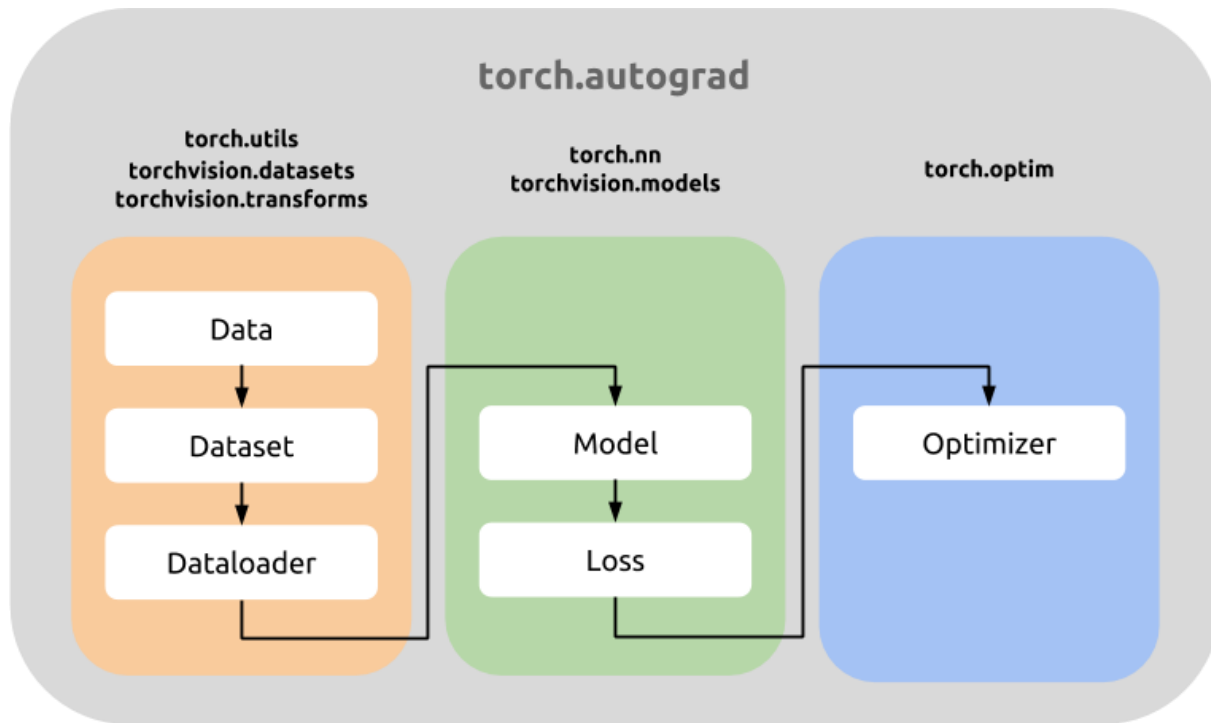


Different Paradigms



	Tensorflow	Pytorch
Graph Creation	Static/Eager	Dynamic/On Runtime
Similar to	C	Python

Pytorch: Overview



Características clave

- Simple device management

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
print(device)
```

```
print(f"Original device: {x.device}") # "cpu", integer
```


```
tensor = x.to(device)
```

```
print(f"Current device: {x.device}") # "cpu" or "cuda", double
```

cpu

Original device: cpu

Current device: cpu



- Implementación de:
 - Optimizers, etc.
 - Datasets
 - Automatic gradients

airplane


automobile

bird

cat

deer

dog





Creación de redes fácilmente

```
import torch.nn as nn
# defining the model
class Net(nn.Module):
    def __init__(self, input_size=1*28*28, output_size=100):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, output_size)

    def forward(self, x):
        x = self.fc1(x)
        return x

net = Net()
net = net.to(device)
```

**Donde está el
backward pass?**

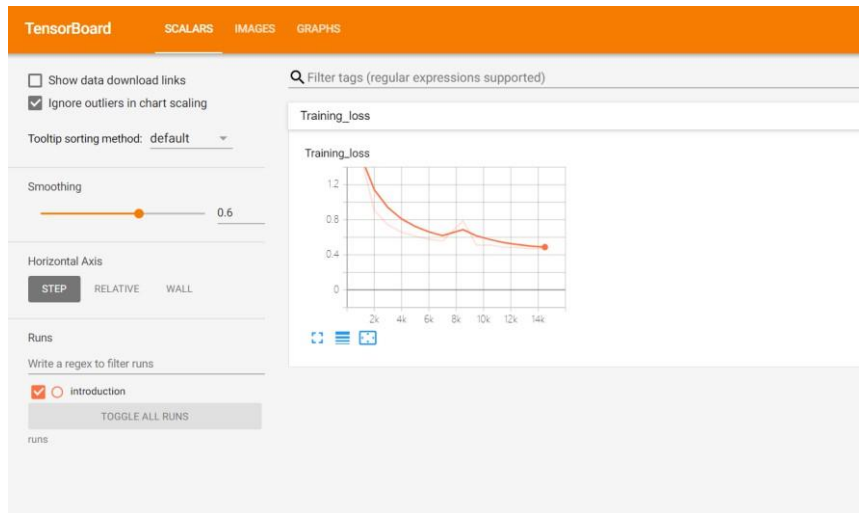


Referencias de Pytorch

- Repository: <https://github.com/pytorch/pytorch>
- Examples (**recommendation**):
<https://github.com/pytorch/examples>
- PyTorch for NumPy users:
<https://github.com/wkentaro/pytorch-for-numpy-users>
- Look up your own and share! 😊

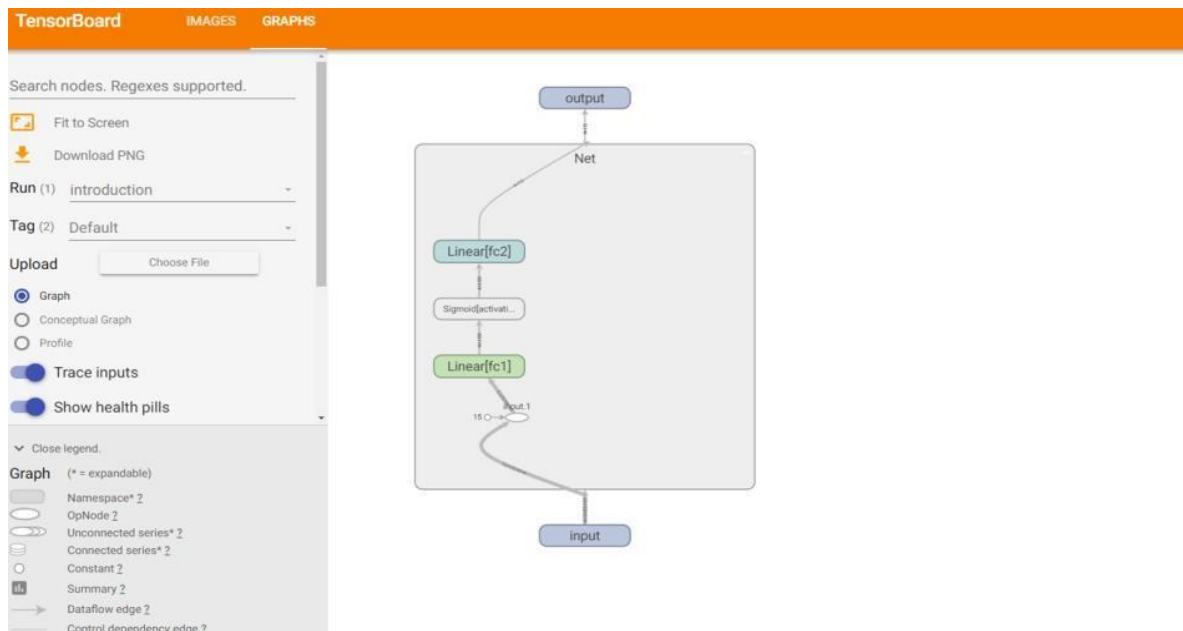
Tensorboard (integrable con Pytorch)

- Directamente accedes a tensorboard en tu training loop
- Tensorboard genera los graph/timestamps etc. a tu gusto

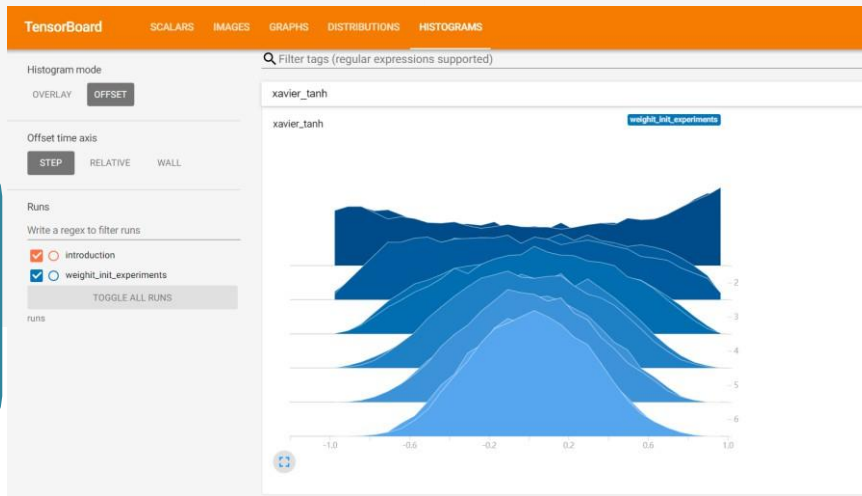
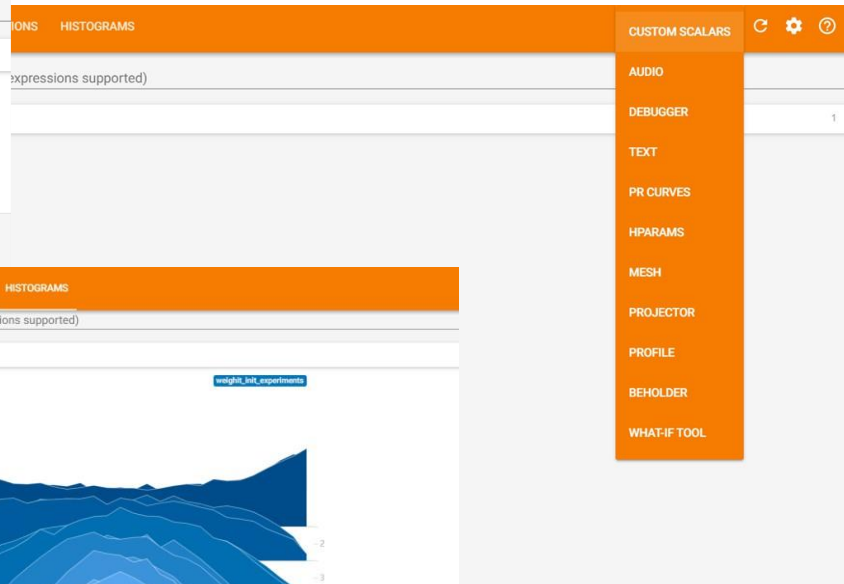


Visualize Networks

- Con un solo forward pass, tensorboard puede mapear y visualizar el gráfico de tu red



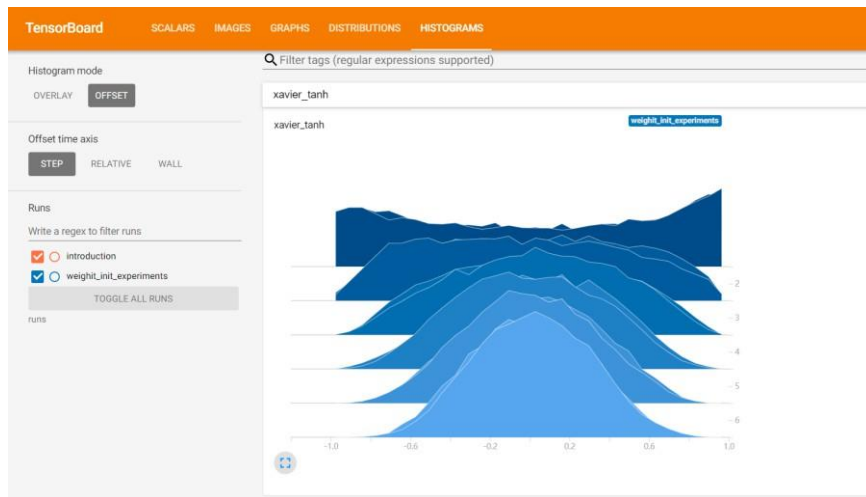

En resumen: ¡documenta todo!



Everything is
saved
“automatically”!

Ejemplo: Weight Initialization

- La visualización del histograma para los layer outputs puede mostrar los efectos de la weight initialization, como se muestra en la lección.



Más abstracción: Pytorch Lightning

Clasificar nuestro código en tres categorías

1. Research code (¡la parte emocionante!, cambia con nuevas tareas, modelos etc.)

→ LightningModule

1. Engineering code (lo mismo para todos los proyectos y modelos)

→ Trainer

1. Non-essential code (logging, organización de ejecuciones)

→ Callbacks

```

# models
encoder = nn.Sequential(nn.Linear(28 * 28, 64), nn.ReLU(), nn.Linear(64, 3))
decoder = nn.Sequential(nn.Linear(3, 64), nn.ReLU(), nn.Linear(64, 28 * 28))

encoder.cuda(0)
decoder.cuda(0)

# download on rank 0 only
if global_rank == 0:
    mnist_train = MNIST(os.getcwd(), train=True, download=True)

# split dataset
transform=transforms.Compose([transforms.ToTensor(),
                              transforms.Normalize(0.5, 0.5)])
mnist_train = MNIST(os.getcwd(), train=True, download=True, transform=transform)

# train (55,000 images), val split (5,000 images)
mnist_train, mnist_val = random_split(mnist_train, [55000, 5000])

# The dataloaders handle shuffling, batching, etc...
mnist_train = DataLoader(mnist_train, batch_size=64)
mnist_val = DataLoader(mnist_val, batch_size=64)

# optimizer
params = [encoder.parameters(), decoder.parameters()]
optimizer = torch.optim.Adam(params, lr=1e-3)

# TRAIN LOOP
model.train()
num_epochs = 1
for epoch in range(num_epochs):
    for train_batch in mnist_train:
        x, y = train_batch
        x = x.cuda(0)
        x = x.view(x.size(0), -1)
        z = encoder(x)
        x_hat = decoder(z)
        loss = F.mse_loss(x_hat, x)
        print('train loss: ', loss.item())

        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

# EVAL LOOP
model.eval()
with torch.no_grad():
    val_loss = []
    for val_batch in mnist_val:
        x, y = val_batch
        x = x.cuda(0)
        x = x.view(x.size(0), -1)
        z = encoder(x)
        x_hat = decoder(z)
        loss = F.mse_loss(x_hat, x)
        val_loss.append(loss)
    val_loss = torch.mean(torch.tensor(val_loss))
model.train()

```

Turn PyTorch into Lightning

Lightning is just plain PyTorch



Submission Opcional

CIFAR10... Again...

- Tarea: Clasificación CIFAR10 (pero ahora en Pytorch)

- Nuevo:
 - Más conocimientos de la clase 7
 - Se puede utilizar todo, pero no: convolutional layers/transformers/pre-trained networks

airplane


automobile

bird

cat

deer

dog



Nos vemos el próximo lunes 😊